

How Generous You Already Are Predicts Whether Fairness Gives or Takes

Dheirav

Responsible AI

Algorithmic Bias Mitigation

dheirav2005@gmail.com

Abstract—A group-fairness constraint requires two groups’ selection rates to be close, but it does not say how the gap should close, and the certifying metric is content with either answer: it reads 0.018 where a demographic-parity constraint destroyed a fifth of the favourable decisions and 0.010 where it created more than it destroyed. A team cannot tell from its own dashboard which it has just done. We show the direction is set by the *population* rather than by which in-processing method is chosen, and is predictable before deployment from the rate at which the current model says yes, read against that population’s own crossover — the rate a team already has, the crossover it must measure. Across 161 disjoint population samples spanning six decision domains, eight data sources and five countries, the constraint withdraws below a population’s crossover and extends above it: on UCI Adult five in-processing methods each shrink the pool by 7.9–22.1%, while on mortgage lending the identical constraint grows it. Committed before ten never-measured populations, the rule scored 9 of 10 against a best constant of 6.

The claim is narrow and sealed tests made it so. Direction is predictable within a population; magnitude is not, and below one percentage point the direction call is a coin flip. The practical output is an audit that locates a team’s own crossover and *refuses* when it cannot: on fifty populations drawn at random, frame and seed fixed in advance, 78% admit a directional rule and 22% do not, and the procedure returns the latter as unanswerable rather than scoring them. Two scope limits belong here: the relationship disappears under post-processing, and in five of the eight sources the favourable decisions are predicted labels rather than allocations actually made. Of nine sealed or registered direction tests, one passed; we report the failures because each removed a stronger claim than the one that survives.

Index Terms—algorithmic fairness, levelling down, demographic parity, in-processing, selection rate, disparate impact

I. INTRODUCTION

That a fairness constraint is silent about *how* the gap closes is well established. Mittelstadt *et al.* [11] made levelling down the centre of a well-known critique; Maheshwari *et al.* [13] observe that it “often goes unnoticed in the overall performance of the model”; Krco *et al.* [12] built an auditing framework precisely because an aggregate score reports neither how many people were affected nor in which direction; and Hu and Chen [6] show, on the dataset this paper opens with, that a parity-proxy constraint can leave both groups with fewer favourable classifications at once. We take all of that as given.

What the literature leaves open is the question a deploying team actually faces: **when** does each direction happen? The question is not academic, because the certifying number is

not merely silent about the answer — it is equally content with either. On UCI Adult a demographic-parity constraint destroyed roughly a fifth of the favourable decisions, and the parity difference it was certified against read **0.018**. On HMDA mortgage data the identical constraint created more favourable decisions than it destroyed, and the parity difference read **0.010**. Both pass. A team looking only at its fairness dashboard cannot tell which of the two it has just done.

The claim, in one sentence. Whether a parity constraint hands out more favourable decisions or takes them away is set by the *population* being decided about rather than by which in-processing method is chosen — and a team can find out which, before deploying, by locating its own crossover and reading its current selection rate against it.

Two words in that sentence are doing work and we would rather name them than have them found. *In-processing*: under post-processing the relationship is absent, so the direction is not a population property in general, only within the family of methods studied here. *Locating*: the selection rate is a number every deploying organisation already has, and the crossover is not — it costs a sweep. The only version of this claim testable from a dashboard alone is a transported prior, whose record is 9 of 10 and then 5 of 10, and which we demote accordingly.

The audit that follows from it has *three* outcomes rather than two — the constraint will give, it will take, or the population admits no answer and the audit says so (Algorithm 1). The third is not a caveat attached to the method; it is a third of what the method does.

Four conditions bound the sentence above. Each is measured rather than assumed, and each was bought with a test that failed:

- **Direction only, and only above a floor.** A sealed model of magnitude lost to predicting zero (MAE 4.50 against 0.77). Effects below one percentage point are called correctly 61% of the time — a coin flip — against 95% above it.
- **About four populations in five.** On fifty populations drawn at random, frame and seed fixed before any arm ran, 78% admit a directional rule and 22% do not. The procedure *refuses* the 22% rather than guessing at them.
- **In-processing, not post-processing.** Under group-threshold post-processing [5] the relationship is absent ($r = -0.024$), and a sealed deconfounding test locates

that boundary at the optimizer family rather than at access to the protected attribute (Section VI).

- **Mostly predicted labels, not allocations.** In five of the eight data sources the favourable decisions are model predictions rather than decisions actually made. HMDA, the one source recording real approvals, supports the direction claim on its natural arms but is also where our sweep protocol failed its held-out test.

The rule underneath is not surprising once stated: say yes often and the constraint tends to give, say yes rarely and it tends to take. What is *not* available without measuring is where the boundary falls for a particular population. Located crossovers span 0.22 to 0.81. In the one setting where we drew populations at random rather than choosing them — US state-level income prediction, fifty draws — no population’s operating rate reached 0.55, and 25 of the 47 with a computable rate sat inside the 0.28–0.65 interval where most of them fall, which is precisely where an unaided reading of “often” against “rarely” has nothing to say. A further fifth of them follow no directional rule at all; there the intuition still returns a confident answer, and the audit declines to.

A. Scope, stated first

This is a paper about **attribute-blind in-processing** constraints, meaning the regime in which the protected attribute is unavailable at decision time. This is the regime most deployed systems occupy: in the United States, per-group decision thresholds face direct disparate-treatment exposure (*Ricci v. DeStefano* [27]; in credit, ECOA and Regulation B prohibit attribute-based terms outright), and commentators read the doctrine the same way [19], [26]. Two honest qualifications. First, the barrier is jurisdiction-specific, and the two major jurisdictions draw it in different places: in the US even *training-time* use of the attribute — which every in-processing method here involves — is legally unsettled, while in the EU the split runs the other way: *decision-time* use of special-category data faces GDPR Article 9 and direct-discrimination doctrine, but training-time use for bias detection and correction in high-risk systems is expressly permitted, under safeguards, by Article 10(5) of the AI Act — and for race and ethnicity the blind regime is often simply factual, because most member states collect no such attribute at all. Second, because of that split we motivate the regime as the one systems occupy, not as one any law cleanly compels. Under post-processing the effect we report disappears, and a sealed deconfounding test shows the disappearance belongs to the method, not to the attribute: the same in-processing optimizer, given the attribute, preserves the relationship (Section VI). The blind regime remains the paper’s scope because it is where the record was built and where deployed systems sit — not because blindness is what the finding depends on.

B. Contributions

The phenomenon. An attribute-blind parity constraint with-draws favourable decisions on some populations and extends them on others, and the certifying metric reports the same

kind of number in both cases: on UCI Adult five in-processing methods each shrink the pool by 7.9–22.1% at an exchange rate of 2.68 destroyed per created, while on HMDA mortgage lending the identical constraint grows it by 4.3% at 0.50. Measured over **161 disjoint population samples**, six domains, eight sources, five countries. Counted in people rather than rates the burden is more lopsided still, and that divergence is itself predictable ($r = +0.885$).

The audit, and its refusal. The same design lets a practitioner locate their own crossover from a deployed model, and return NO ANSWER when the population supports none. On fifty populations drawn at random — frame, seed and draw committed before any arm ran — 78% admit a directional rule and 22% do not, and the procedure hands back the latter rather than scoring them. Detecting a non-monotone landscape requires the sweep that transporting a prior exists to avoid, which is why the refusal cannot be replaced by a dashboard reading.

The predictor, identified by elimination rather than asserted. Holding the task fixed and moving only the decision rule separates the rate from task difficulty. A screen-gated Brazilian race cohort — ten arms from two person-samples, so the confound breaks *within* populations — beat a cutoff-only null 8–9 to 5; a sealed lending cohort beat a loan-purpose null 8 to 5 on real approvals; a five-domain comparison rules out the group gap. We say *predicts*, never *causes*: no pooled slope transfers.

A conditionality theory predicts, now measured. Contemporaneous theory [14] establishes distribution-dependence in the blind regime without experiments, its conditions being orderings on score regions rather than measurable quantities. A percentile-relaxed form of that ordering agrees with the measured direction on 24 of 26 arms and tracks our predictor at $r = +0.935$, holding across 150 arms at $r = +0.850$ (population-clustered 95% interval $+0.813$ to $+0.892$) and surviving a boosted-tree probe. The exact conditions were certified on none of the 26 by our plug-in estimator, whose ratio diverges in the $\hat{\nu}$ tail — an estimator pathology, and we tried no other.

Two predictive claims, and a prior demoted. A *sweep-conditional* claim — measure your own crossover, then read your rate against it — and a weaker *transported-prior* claim, the only one testable before any measurement. The rule itself is sealed: committed before ten never-measured populations, it scored 9 of 10 against a best constant of 6. The fixed *value* 0.54 fares worse and we no longer present it as a finding: located crossovers span 0.22 to 0.81, the value failed its own strict bar on the race cohort, and on a second post-hoc cohort it scored 5 of 10 — each miss agreeing with that population’s own sweep.

Boundaries stated rather than left for review to find, including a computable test for when the procedure cannot be used at all. Of nine sealed or registered direction tests, one passed; each failure removed a stronger claim than the one that survives, and one cohort built to strengthen the paired statistic showed the attempt self-defeating by construction. The

relationship survives changes of route, tolerance, base learner and protected attribute, and disappears under post-processing.

II. RELATED WORK

What we take as given. Levelling down as a phenomenon and minimum-rate constraints as a remedy are Mittelstadt *et al.*'s [11]; the impact decomposition we use throughout is Krco *et al.*'s [12]; fairness gerrymandering is Kearns *et al.*'s [15]; the income-cutoff knob is Ding *et al.*'s [16]; the reductions we constrain with are [1]. Contemporaneous theory [14] establishes that in the blind regime the direction is distribution-dependent, so we do not claim the conditionality itself — what it contains is no experiments and conditions stated as orderings on score regions rather than measurable quantities.

Four adjacent results, and what each leaves open. Hu and Chen [6] come closest to the phenomenon: an in-processing parity-proxy constraint on *this paper's own dataset* leaving both groups with fewer favourable classifications, read as a Pareto violation. Their existence claim is prior to ours and we do not repeat it as new; what they do not do is parameterise the effect, and their conclusion — that welfare moves non-monotonically in the fairness tolerance within one population — concerns a different axis from ours, which compares unconstrained against constrained across populations. Liu *et al.* [7] is the closest structural analogue: a critical selection rate at which the sign of a parity constraint's effect flips, with the baseline's position relative to it doing the predicting. Three things differ and the resemblance is worth stating plainly rather than leaving for a reader to find — their flipping quantity is one group's expected score change over a time step against our one-shot headcount over both groups, their axis is the protected group's own rate against our model's overall rate, and their policies are group-specific where ours are blind. Menon and Williamson [8] give the parity-optimal blind classifier as a per-instance threshold correction, pushing some subjects out of the positive set and pulling others in without evaluating the net. The infra-marginal line [9], [19] supplies the vocabulary a mechanism would need — what happens at the margin, not away from it — and a direction flip of its own, though theirs is one group's share of a fixed budget rather than a pool that can change size.

On not explaining the mechanism. Our sealed derivation lost to a constant (Section IX) and we claim no mechanism, but we should not overstate that into an absence of any account. For the attribute-aware Bayes-optimal case the parity-optimal thresholds are available in closed form [10] as opposite-signed shifts each inversely proportional to a group's size; to first order the group sizes cancel and the sign of the pool change is set by which group carries more score density at the decision margin — an infra-marginality statement, and a property of the population rather than of the method, which is what we measure. We could not extend it: our constraint is blind, our classifier is a reduction's randomized mixture rather than a threshold rule, and the effects we measure are far outside a first-order regime. The honest statement is that we failed to

extend a known marginal argument to the blind randomized case, not that no account exists.

III. SETUP AND METHOD

Data. 161 independent populations across eight data sources. Two are US surveys: UCI Adult (45,222 rows) and ACS Income [16] across 102 state-years and two protected attributes. One records real allocations: HMDA 2018 mortgage filings, 66 rate-bearing arms over 40 states (30,000–305,593 records each). The remaining five are IPUMS International census extracts for Brazil and Mexico, COMPAS [18], LSAC bar passage, the Dutch national census of 2001 and Taiwanese credit-card default records of 2005. A *population* is one such unit — a state, cohort or census — and *independent* means disjoint person samples: no person appears in two, and arms within a population share people, so however many arms a population contributes it is counted once. Ratios here are over populations unless they say otherwise, because four counts in this project's own history turned out to be arm counts wearing population clothes.

What a favourable decision is, and mostly is not. This belongs here rather than in Limitations, because it bounds everything that follows. **In five of the eight sources no one receives anything:** the favourable decisions are a classifier's labels over people whose incomes, occupations or bar results already exist, so the pool being given and taken is a pool of predictions. HMDA is the one source recording real allocative decisions — and it is where the direction reverses, and where our sweep protocol failed its held-out test. Whether the construct transports from status-prediction benchmarks to allocation data is an open question rather than an assumption we are entitled to.

The lending outcome, defined. On HMDA the outcome is the *lender's* approval: action-taken codes 1 (originated) and 2 (approved, not accepted) against 3 (denied), with withdrawn, incomplete, purchased and preapproval records excluded as not being a completed approve-or-deny. Features are restricted by allow-list to what exists when the application is filed. The race arm is White vs. Black applicants, the sex arm Male vs. Female (joint applications excluded), and loan purposes are sub-populations.

COMPAS, LSAC and the Dutch census were added to leave the sources the finding was formed in. The Dutch census carries a between-group gap of 0.298, roughly twice Adult's, making it the test of the standard alternative explanation. LSAC was chosen because bar passage is naturally generous (base rate 0.890).

Definitions, in one place. An *arm* is one population under one configuration (a decision threshold, a label cutoff, a tolerance, a learner). The *selection rate* r is the fraction of the test split the baseline model marks favourable. The *pool change* Δ_{pool} is the percentage change, baseline to mitigated, in the number of favourable decisions; *up/extension* and *down/withdrawal* name its sign. A population's *crossover* is the rate at which its pool change changes sign, reported as the bracket between the highest rate that still shrinks the pool

and the lowest that grows it. The *viable band* is the span of selection rates over which the baseline stays above the trivial predictor. The *exchange rate* is favourable decisions destroyed per favourable decision created, an unweighted descriptive count.

Protocol. The baseline learner is an L2-regularised logistic regression on one-hot-encoded categorical and standardised numeric features; one section swaps in a boosted tree and says so. Each arm draws a 70/30 train/test split, stratified on group $\times$ label, re-drawn per seed. Five seeds per arm unless stated (the replication checks use twenty); every per-arm quantity in this paper is a mean over its arm’s seeds, and no single-seed number is comparable to the tables.

Convention. $y = 1$ is always the favourable outcome. This required inverting COMPAS’s recorded target, which encodes recidivism, and declaring *Female* the advantaged group on its sex arm, since women in that cohort reoffend less. Both are enforced by automated checks, because either error produces every directional result sign-inverted without raising.

Mitigation. `ExponentiatedGradient` [1] at $\varepsilon = 0.01$ unless stated (ε is the library default and bounds the *training-time* empirical moment violation of the randomized classifier, so held-out gaps such as Table II’s 0.018 can lawfully exceed it; the tolerance sweep spans 0.002–0.05, a $25\times$ range). Fairlearn 0.14.0 throughout, default `ExponentiatedGradient` parameters beyond ε . No model reads the protected attribute at prediction time except where post-processing is explicitly under test.

Exclusion rules. An arm is discarded if the baseline demographic-parity gap is below 0.05, or if the baseline classifier’s accuracy falls below $\max(p, 1-p)$ — the accuracy of always predicting the more common label. The second rule matters more than it sounds (Section IX). Provenance, because exclusion rules can select on outcomes: both rules were developed *during* the exploratory phase, in response to failures the ledger records, so retained-arm correlations from that phase are conditioned on in-sample-tuned filters and the sensitivity table reports the dependence. What the sealed tests then did with the two rules varies, and stating it as uniform would flatter them: the re-seal applies *neither* and scores all ten populations raw; the third-direction and lending cohorts apply the magnitude guard their own protocols fix, but not these two; and the screen-gated race cohort is the one that carries both, frozen before its extract existed. The sealed record is therefore less filtered than an unqualified claim of frozen guards would suggest, which is the direction that favours it.

Accounting. Table I states which populations enter which analysis; the complete record gives, per headline ratio, the arms, the disjoint person samples and the instruments it is counted over. Four counts in this project’s own history turned out to be arm counts masquerading as population counts, so a ratio here is a ratio over populations unless it says otherwise. The second table is computed from the stored results by `scripts/independence.py` rather than maintained by hand, since the counts it reports are exactly the kind that drift. Read together they make one distinction unavoidable: *a ratio*

TABLE I
WHICH POPULATIONS ENTER WHICH ANALYSIS. AN *arm* IS ONE POPULATION UNDER ONE CONFIGURATION; POPULATIONS ARE COUNTED ONCE HOWEVER MANY ARMS THEY CONTRIBUTE.

Analysis	Populations (arms)
Ablation; who pays	Adult (5 methods $\times$ 5 seeds)
Rate-people divergence	10 pops. (19 arms, 2 attributes)
Label route (income cutoffs)	AL sex; AL, MS, OR, UT race (4 each)
Natural transition	HMDA MS+LA (5 loan purposes)
Threshold route, dense sweeps	Dutch, SC, COMPAS, AL, KY (12 pts.)
Tolerance, $25\times$ range	AL, CT, KY, OR, SC
Boosted trees	the same five states
Intersectional	Adult + 10 populations
Post-processing regime	17 populations ($r = -0.024$)
Attribute-aware direction	27 of 27, of which 9 pre-registered
Relaxed ζ ordering	26 arms from 15 populations
Crossover intervals	4 non-lending + 3 lending estimates
Sealed test 1 (failed)	8 scored + Taiwan, all fresh
Re-seal (holds)	10 fresh ACS states
Residual campaign (underpow.)	10 fresh states (76 arms)
Sealed magnitude (failed)	NY + TX (16 arms)
Cross-year shape seal (failed)	8 fresh state-years (70 runs)
Attr.-independence seal (failed)	6 race arms (48 runs)
Selection-rate floor	10 pops. (19 arms)
Cross-task shape seal (failed)	6 task arms, 4 states’ persons; not indep.
Third direction cohort (failed)	24 fresh state-years, 2014 + 2019
Sealed lending cohort (failed)	8 fresh markets (16 arms, 2 purposes each)
Lending coverage	40 race + 26 sex arms (66)
Landscape survey	50 pops. at random (660 arms)

TABLE II
FIVE IN-PROCESSING METHODS ON UCI ADULT, FIVE SEEDS: THE TWO REDUCTIONS AND GRIDSEARCH FROM [1], ADVERSARIAL DEBIASING AFTER [2], AND THE PREJUDICE REMOVER AFTER [3]; DI IS THE DISPARATE-IMPACT RATIO OF [4]. EVERY ONE REACHES PARITY, AND EVERY ONE SHRINKS THE POOL.

Method	Acc.	DP	EO	DI	Δ pool
Unconstrained	0.847	0.186	0.095	0.300	—
ExpGrad (DP)	0.828	0.018	0.280	0.895	−20.5%
GridSearch (DP)	0.827	0.015	0.304	0.986	−22.1%
Adversarial deb.	0.826	0.020	0.250	0.889	−14.8%
Prejudice remover	0.837	0.065	0.188	0.686	−10.2%
ExpGrad (EO)	0.836	0.107	0.032	0.521	−7.9%

over arms is not a ratio over populations, and where the two differ this paper says which it means.

IV. PARITY IS REACHED BY WITHDRAWAL — IN SOME POPULATIONS

Table II gives the ablation. Every method reaches its target: demographic parity falls from 0.186 to as little as 0.015 for under two accuracy points, and this holds on a linear model and on an adversarial one alike. The mitigation itself works, which is why we treat it as background rather than as a contribution.

What changes alongside is the part the metric does not report: **every one of these methods shrinks the pool**, by 7.9% to 22.1%, and not one of them closes the gap primarily by extending decisions to the disadvantaged group. Eighteen of nineteen survey arms show the same pattern, which suggests

TABLE III

WHO PAYS, UCI ADULT. THE RATE-LEVEL VIEW AND THE PEOPLE-LEVEL VIEW OF THE SAME INTERVENTION DISAGREE SYSTEMATICALLY.

Method	Share (rates)	Share (people)	Exchange
ExpGrad (DP)	0.575	0.742	2.68
GridSearch (DP)	0.575	0.743	2.69
Adversarial deb.	0.508	0.700	1.96
Prejudice remover	0.498	0.673	2.04
ExpGrad (EO)	0.531	0.660	1.92

this is a property of the populations rather than of any one method.

A. Rates understate what is happening to people

The decomposition of [12] splits the closed gap into the part paid by the advantaged group losing ground and the part gained by the disadvantaged group. Measured in *rates*, that split is close to even — 0.50 to 0.58 of the closure comes from withdrawal (Table III).

Measured in *people* the split is not even. Between 0.66 and 0.74 of the individuals moved in the two directions that close the gap — advantaged subjects losing a favourable decision, plus disadvantaged subjects gaining one — are advantaged subjects losing one. The denominator is those two flows and not every changed decision: it excludes the smaller counter-flows in each group, and a reader who takes it for all churn will find it inconsistent with the exchange rate, which is counted over every flow. The reason the two views disagree is that equal-looking rate changes fall on groups of unequal size, so the rate-level view — which is what a fairness dashboard reports — understates the burden counted in individuals, and it does so in every method tested. The divergence between the two views is itself predictable from a cross-flow diagnostic, at $r = +0.885$ across populations.

The **exchange rate** shows the same thing without needing a share: ExpGrad under demographic parity destroys 2.68 favourable decisions for every one it creates.

The normative baseline, stated as the convention it is. Throughout, “withdrawal”, “who pays” and the exchange rate are counterfactual-comparative quantities measured against the *unconstrained model’s* allocation — a baseline with no privileged moral standing, since the constraint exists precisely because that allocation is suspect. On a claims-based view [28], withdrawing an approval that answered to no legitimate claim wrongs no one; a positive classification is also not automatically a benefit (extension in lending can be predatory inclusion). The exchange rate is therefore an unweighted descriptive count, not a welfare ledger: a reader with prioritarian weights on the disadvantaged group’s gains may score a rate above 1.0 as net-good, and the group-conditional flows the tables report make such weighted readings computable. What the accounting establishes is narrower and metric-relative: the certifying number cannot tell these outcomes apart, whatever one’s weights.

On HMDA mortgage data, however, the direction reverses. The identical constraint *grows* the pool by 4.3%, at an exchange rate of 0.50 — two favourable decisions created for every one destroyed. What makes this worth noticing is that the parity metric cannot see the difference: it reports 0.018 on the population that destroyed a fifth of its favourable decisions, and 0.010 on the one that created more than it destroyed.

V. WHAT PREDICTS THE DIRECTION

A. A single-factor design

ACS Income’s label is “earns more than \$50,000”, and that threshold is a benchmark convention rather than a fact about the world, which makes it a knob we can turn. Moving it on one fixed state varies the selection rate while the sampled people, instrument, feature set and group ratio stay fixed — though moving the label redefines the prediction problem itself, so this is one arm of the two-routes triangulation below, not an isolation of the rate on its own. The change in the pool rises with the baseline selection rate at $r = +0.801$ over four arms — too few points for formal inference, and reported as descriptive; a partial correlation on four points has one degree of freedom and is omitted as uninformative. The sign flip is the substantive observation: 22 decisions are destroyed per one created at a rate of 0.03, compared with 0.75 at 0.89.

B. The transition appears without being manufactured

The same transition also appears where nothing is manipulated at all. Pooling two states across five real loan products, home improvement at a selection rate of 0.555 levels *down* (−1.57%) while refinancing at 0.871 levels *up* (+2.95%), with $r = +0.803$ and Spearman $\rho = +0.900$. This means a bank running one constraint across its loan book would take opportunities away in one product while handing them out in another, and its fairness report would show success in both.

C. It is the rate, not the difficulty of the task

The design above moves the *label*. “Earns over \$100,000” is not a rarer version of “earns over \$20,000”; it is a harder problem. That design therefore cannot separate two explanations which predict everything observed so far.

We separate them by holding the task completely fixed — same rows, features, groups and *fitted scores* — and moving only where the model draws its line. The relationship reproduces, and so does the crossover’s location: on Oregon the two routes give 0.362–0.653 and 0.353–0.637 (Fig. 2). **Task difficulty is not doing the work.**

With twelve points chosen from each population’s viable band, the relationship holds on the Dutch census (+0.946), South Carolina (+0.905) and COMPAS (+0.844). Alabama and Kentucky return *negative* values on eleven and ten arms; their viable bands top out at 0.566 and 0.561, at or below every crossover measured, so both sweeps lie almost entirely on one side of the transition. **We therefore claim the relationship across the crossover and withdraw any claim that it holds within the low-rate region alone.**

In 6 of 8 populations the change in the pool rises with the selection rate and crosses zero. The 2 that reverse are marked. Grey = selection rates no deployable classifier can reach on that population, which cuts some off above the crossover and others below it. Red band = the 0.51–0.58 cluster.




Fig. 1. Every retained arm from every densely swept population. Within a population the change in the pool rises with the baseline selection rate and crosses zero; the two that reverse are marked, and both are sweeps whose viable band confines them below the crossover. Panels are deliberately *not* pooled: no pooled slope survives its pre-registered bar, so one scatter would assert a common curve the data reject. Vertical band: the point-estimate cluster of Table VI, drawn as first published; Section VIII gives its later history.




Fig. 2. The design that separates the selection rate from the difficulty of the task. Moving the income cutoff changes the label, and therefore how hard the problem is; moving the decision threshold changes neither, only where the fitted model draws its line. Both routes cross zero inside the same band. If difficulty were doing the work, the second curve would not reproduce the first.

D. Five domains, and the group-gap alternative

Fig. 1 shows every arm, and Table IV the summary. The pre-registered alternative — “this is a property of income pre-

diction and American mortgage lending”, requiring $|r| < 0.30$ — is refuted on both new instruments, and refuted on their

TABLE IV

THE RELATIONSHIP ACROSS INSTRUMENTS, NOT POOLED ACROSS THEM, FOR THE REASON FIG. 1 DOES NOT POOL. THIS TABLE’S JOB IS *breadth* — ONE POPULATION PER INSTRUMENT, ON WHATEVER ARMS THAT POPULATION NATURALLY SUPPLIES — SO THE FOUR-TO-SIX-ARM CORRELATIONS ARE DESCRIPTIVE, EXACTLY AS SECTION V-C CONCEDES FOR ALABAMA. THE DENSE-SWEEP COLUMNS CARRY THE WEIGHT WHERE A VIABLE BAND ALLOWS TWELVE POINTS, BUT THOSE TWELVE ARE THRESHOLDS APPLIED TO ONE FITTED SCORE VECTOR ON ONE SAMPLE AND ARE NOT TWELVE INDEPENDENT OBSERVATIONS, SO WE QUOTE NO p -VALUE ON THEM: A PEARSON TEST THERE DESCRIBES WHETHER A SWEEP IS MONOTONE, NOT WHETHER THE RATE PREDICTS ACROSS POPULATIONS. THE PRE-REGISTERED ALTERNATIVE IS REFUTED ON THE DENSE COLUMNS, NOT THE SPARSE ONES.

Instrument	Domain	Population	natural arms		dense sweep	
			r	n	r	n
ACS Income	income	Alabama	+0.801	4	—	—
ACS Income	income	S. Carolina	—	—	+0.905	12
HMDA	mortgages	MS+LA pooled	+0.858	4	—	—
COMPAS	crim. justice	race arm	+0.870	5	+0.844	12
Dutch census	occupation	sex arm	+0.915	6	+0.946	12
LSAC	legal educ.	—	unsweepable	—	—	—

densely swept arms rather than their sparse ones: COMPAS gives +0.844 and the Dutch census +0.946 over twelve points each. Earlier versions of this paper attached $p < 0.001$ to those coefficients; we withdraw it. Twelve thresholds on one fitted score vector over one sample are near-deterministic transformations of each other, so the Pearson null is the wrong reference and the effective sample size for a claim about *populations* is one. The population-level statistic this design supports is much weaker: across the six densely swept populations the within-population slope is positive in four and negative in two, a sign test at $p = 0.34$. We report it because it is the honest number, and note that the two negatives are the sweeps Section V-C withdraws for sitting entirely below their crossover — which is a stated exclusion, not a rescue, and either way the prospective evidence for this paper lives in the sealed forecasts and not in these coefficients. The four- and five-arm rows are reported for breadth and carry no inferential weight on their own.

The oldest competing account is that the *gap between groups* drives the direction. The Dutch census has a gap of 0.298, roughly twice Adult’s, making it the population where that explanation should win. It gives $r = +0.915$ across all six arms, none excluded.

Because the parity-gap floor was tuned in-sample (Section III), the *natural-arm* column of Table IV was re-scored at floors 0.02/0.05/0.08/0.10: each of those correlations stays between +0.80 and +0.92 with no sign change at any floor (HMDA is floor-invariant outright; Alabama thins to two arms only at 0.10). **The dense-sweep column is a different matter, and an earlier version of this paper claimed otherwise.** Table VII’s 0.05 column *is* that dense-sweep column, and one of its three entries is strongly floor-dependent: South Carolina reads +0.905 at the committed floor and +0.012 at 0.02 and 0.01. So the floor is doing real work on a number this paper leans on, and the correct statement is the narrower one — the natural-arm evidence is floor-robust, the dense-sweep

TABLE V
WHAT EACH MANIPULATION PRESERVES

Manipulation	Direction	Magnitude
Route to the rate	unchanged	differs 3×
Tolerance, 25× range	unchanged	differs 4×
Boosted trees (5/5 pops.)	unchanged	differs 3×
Attribute, sex → race	unchanged	—
Criterion → equalized odds	weaker, unevenly	differs 8×
Optimiser → post-proc.	vanishes	—

evidence is not uniformly so, and South Carolina’s twelve-point correlation should be read as conditional on the exclusion rule rather than as independent of it.

E. Robustness

Table V summarises. **Within a population, and in the attribute-blind regime, the selection rate predicts which way and orders how much. Neither transfers between populations:** the signed distance from a population’s own crossover gives ρ of +0.96, +0.95 and +0.78 in three of four populations, but a pooled slope reaches only $r = +0.487$ against a pre-registered +0.70 bar.

VI. THE REGIME BOUNDARY

The relationship stops where the optimizer family changes, not where the protected attribute becomes readable, and a sealed test with both readings’ bars committed in advance is what separates those. Under group-threshold post-processing [5], at each population’s natural operating point, the rate carries no information about the pool’s direction ($r = -0.024$, against +0.585 for the reduction on the same populations); our pre-registered prediction that the rule would carry over is a recorded failure. The missing cell decides between the two available explanations: given the attribute under the *same* in-processing optimizer, the relationship survives at $r = +0.672$ across thirteen natural arms — the method reading, not the blindness reading. Two cautions we would rather state than have computed for us: at thirteen arms the interval on that coefficient is wide enough to touch the competing reading’s region, and its difference from +0.585 is not itself a difference. What carries more weight than either coefficient is the mechanism the failure exposed — the post-processor re-derives its own thresholds and returned an identical model at all six operating points, so the measured effect was a constant divided by a shrinking baseline.

The theory’s group-level claim is untouched. Where the attribute is readable the advantaged group loses and the disadvantaged gains, without exception across our attribute-aware arms. That is a consistency check against a published theorem rather than a forecast — the theorem made the prior close to one — and it separates cleanly from the pool question: the theorem fixes each *group’s* direction wherever the attribute is readable, while the *pool’s* direction, the sum it does not constrain, stays rate-predictable under in-processing in both regimes and unpredictable only under post-processing.

A percentile-relaxed form of the theory’s ordering agrees with the measured direction on 24 of 26 arms and tracks our predictor at $r = +0.935$, so the selection rate is a cheap proxy for a quantity the theory states expensively; the exact conditions were certified on none of the 26 by our plug-in estimator, whose ratio diverges in the $\hat{\nu}$ tail, and we tried no other.

VII. THE AUDIT: GIVE, TAKE, OR REFUSE

Algorithm 1 states the procedure. Three of its steps are worth drawing out, because each is a place where a plausible-looking run returns a meaningless answer.

It can refuse. Lines 6–8 compute the viable band and exit before any constraint is fitted, if no threshold reaches a low selection rate with a model still beating the trivial predictor. On LSAC that band is too narrow to sweep (Section VIII) and the procedure declines, rather than returning arms that would be discarded later.

It can return VOID. Line 16 requires enough arms to survive the guards and the pool to have moved at all. A correlation computed over arms that barely move is fitted to noise: Connecticut returns $r = -0.924$ on a spread of a tenth of a percentage point, and without this test that number reads as a refutation of the whole finding.

It can return NON-MONOTONE, and on current data it often will. Lines 20–22 exist because the assumption behind bracketing — one rising sign change — is itself population-dependent: among the ten largest states we swept, three are U-shaped and the 2022 vintages at nominal labels invert, so a bracket step that assumed monotonicity would return a confident wrong direction precisely there. A NON-MONOTONE verdict means the directional rule does not apply to this system, which is a first-class answer, not a failure of the audit. A sweep whose retained arms all share one sign is the opposite case and returns that direction (line 24): the constraint moves the pool the same way across the whole viable band, which is what “levels up throughout” means operationally. Two cautions the earlier version of this algorithm kept in prose are now steps: arms below a selection rate of 0.10 are marked *advisory* and excluded from the shape call (line 14), because below that boundary the two routes to a rate give opposite signs (Section VIII) and the accuracy guard does not catch it; and every fitted arm is probed for the flat keep-probability signature (line 19), because a lottery under the certificate is a finding the direction verdict alone would hide. The probe’s decision rule, stated so it can be implemented: flag an arm whose mean granted probability just below the boundary is under 0.10 while the keep-probability’s correlation with the score above it is ≈ 0 — the observed separation is wide (0.00–0.07 against 0.17–0.43, and ≈ 0 against +0.37 to +0.73).

Guard provenance. A guard’s provenance belongs to the pair it forms with a result, not to the guard. The band span, the parity-gap floor and the threshold grid are exploratory, fixed in the threshold sweep’s own pre-registration and confirmatory for every seal after it; three others are not, and the complete record states for each whether it was fixed before, after, or —

in one case, a void guard imported backwards from a later pre-registration that moved three arm sets from refutation to VOID — applied retrospectively in the conjecture’s favour, against which the same audit found no arm set that passed on noise.

It can return INDETERMINATE. Lines 28–29 refuse a sign call when the deployed rate sits inside the located bracket, or when the nearest arms’ effects are within seed noise — precisely the two conditions under which this project’s own record shows a sign call has nothing to grip: Maryland’s miss sat at the cluster’s edge (Section VIII-C), and Minnesota’s was an effect of -0.04% that flips sign across seeds. The registered rubric of the re-seal did not have this verdict and was charged a miss for it; the algorithm inherits the lesson.

The magnitude guard is measured, not assumed. Line 28’s refusal rests on the claim that a near-zero effect has no sign to predict. Two sealed cohorts run after this algorithm was written put a number on it: across their forty arms the rule is correct on **21 of 22** arms whose effect exceeds one point and on **11 of 18** below it — 95% against 61%, which is a coin flip. Those arms span 19 and 15 disjoint person-samples, and clustering on the sample rather than the arm leaves the split intact: 85–100% above the threshold against 37–82% below, intervals that barely meet.

What that measures is the floor’s existence, not its location. The two read alike and are not the same claim. Sweeping the guard over the same forty sealed arms, every value from 0.25 to 5.00 points separates the reliable arms from the unreliable ones, by between 13 and 35 percentage points, with the largest separation at 0.25 and the committed 1.0 giving +34. So the guard’s *existence* is measured and its *value* is operational — a round number fixed in a protocol before its arms existed, which the data neither singles out nor contradicts.

It returns a sign, never a size. The distance from the crossover *orders* the effect within a population, so a team can rank its own segments; but no slope transfers between populations, so lines 32 and 34 return a direction and stop there.

What it is for, and what it is not. The audit applies where the pool of favourable decisions can actually change size. Where the allocation is capacity-pinned — a fixed number of seats, beds, or slots — the total is constant by construction, the directional quantity this paper measures cannot arise, and the audit is out of scope from line 0; this is the fixed-resource regime of [17], and clinical triage and admission commonly live there. No natural operating point below 0.13 has been probed in this record, so deployed rates below that carry an untested-region caution on top of every verdict.

What it needs, and what it costs. The audit requires a held-out sample carrying the protected attribute: audit-time access to a , which is legally and practically distinct from decision-time use and is the regime bias-testing provisions contemplate (Section I-A) — in US credit, the self-test privilege of Regulation B (12 CFR §1002.15) exists for exactly this, though it is conditional: the privilege attaches to a lender’s own voluntary self-test and expects appropriate corrective action,

and an analysis of public HMDA records, like this paper’s, is not itself a privileged self-test. It fits the constraint $12 \times$ seeds times; with five seeds and the linear learners used here, one fit takes 2–10 seconds at 15k–69k training rows (measured on a laptop CPU), so a full audit is minutes per population, and its cost scales linearly in the base learner’s own fit time since each reduction call is a sequence of learner refits. A cheap pre-screen exists: the relaxed- ζ ordering of Section VI needs no constraint fits and correlates with the predictor at $r = +0.935$.

What the sweep buys that one fit does not. A team able to run $12 \times$ seeds constrained fits could instead run *one*, at the operating point it already uses, and read the pool change off it — so what has prediction bought over measurement? A single fit returns a point; the sweep returns a neighbourhood. Two measurements separate them. Across 639 adjacent arm pairs whose effects both clear the 1.0-point guard, 69 — 11% — disagree on the *sign*, so the direction one fit reports is not reliably a property of the population rather than of the threshold that fit happened to use. And populations sit close to their own boundary: of 21 disjoint person samples carrying both a located crossover and a natural arm, 11 sit within 0.10 of the crossover and 6 within 0.05, Florida 2018 sitting 0.004 from it. For those teams the direction is not a fact to accept or refuse but a threshold they already control, and one fit reports the withdrawal without reporting the lever. That the crossover cannot be looked up rather than measured shows in the same set: Florida enters twice, once per attribute, crossing at 0.284 read by sex and 0.439 read by race — the same people, two boundaries.

The limit belongs with the answer. The sweep does not buy a better estimate of what the constraint does at today’s threshold; repeated fits at one point buy that. What it buys is *shape*. Five of the 21 sit more than 0.20 from their crossover and would have decided the same from one fit — and that 24% over-states the waste, because every sample in the set has already passed the monotonicity screen, and the 22% admitting no rule never enter it. A team cannot learn which group it is in without running the sweep, which is the reply in one line: the audit is not circular, because its output is not the quantity a single fit measures.

What it actually returns, totalled. Run over the 52 population-label pairs whose stored sweeps resolve to a dataset spec — fourteen further IPUMS sweeps are stored but unmapped, and are excluded rather than guessed at — the distribution is: 29 directional verdicts (20 WITHDRAWAL, 9 EXTENSION), 8 NON-MONOTONE, 13 VOID, one refusal at the noise floor (COMPAS), and one bracket with no natural arm to read against. Every one of the 29 directional verdicts matches what the fitted constraint did at that population’s natural operating point — a consistency check on the audit’s reading of its own sweep, not independent validation, since the natural arm is part of the curve being read. That distribution is computed over the sweeps this project happened to accumulate, which were never randomly chosen, so it answers a question about this project rather than about populations. A survey with a pre-committed frame answers the latter.

The frame, since “at random” means nothing without one. Every combination of the fifty US states, the 2014, 2018 and 2019 ACS vintages, and the two protected attributes, with the income cutoff held at \$50,000 — 300 population-label pairs. The 2022 vintage is excluded because its inversion is a diagnosed nominal-label artifact, and including it would conflate a label problem with a landscape property. Fifty pairs were drawn uniformly at random under a fixed seed, and the frame, the seed, the draw and the exclusion were all committed to version control before any arm ran; the drawn list is in the repository. The draw came out balanced without being forced: 25 of each attribute, and 14, 14 and 22 across the three vintages. The fifty pairs span **48 disjoint person-samples**, two state-years having been drawn under both attributes.

The result: **78% admit a directional rule and 22% do not**, of which 64% are classic single-crossing landscapes. Resampling *populations* rather than draws puts the directional share at 66–89%, and leaving each population out in turn never moves it below 77%. The convenience sample was mildly optimistic — it implied 15% non-monotone — and the larger number is the one to quote. The refusals are not silences. A VOID arm is one whose pool change spans under two points across the entire viable band — median 1.47 — which tells a deployer that the direction question is *moot* on their system, not that it is unanswerable; a NON-MONOTONE verdict says the rule does not apply here, which is a finding about the system; and the noise-floor refusal says the population is too small to audit at all. Counted that way, 48 of the 52 return something a team can act on and four return nothing — three sweeps too thin to read and one bracket with no natural arm. The honest summary: the audit gives a *direction* a little under six times in ten, tells a team the question is moot or the rule inapplicable in most of the remainder, and is silent on four of fifty-two. Finally, an ex-ante verdict earns trust only if checked: whatever the verdict, including refusals, the deployed system’s measured pool change and who-pays decomposition (Section IV-A) should be reported after the fact — the verdict is falsifiable at deployment, and Section XI makes this a requirement rather than advice.

The direction survives the extraction a deployer must apply. Every direction result above is measured on the randomized mixture, but a lender under adverse-action-notice duties deploys a *deterministic* model, so we tested the extraction directly: thresholding the mixture’s per-person probability at 0.5 on the natural arms of ten populations spanning both directions. The extracted model agrees with the mixture’s pool direction on **10 of 10**, with the parity repair largely intact (extracted gaps 0.009–0.034 against baseline gaps 0.08–0.19). One deployment fact surfaced on the way: an extraction that instead tries to reproduce the mixture’s own operating point cannot generally do so, because the mixture’s probability takes few discrete values and the ties force the deterministic threshold off the target rate — on three cutting populations far enough off to flip the sign of the measured pool change. Thresholding at 0.5 is the well-defined extraction, and it is the one that preserves the verdict. The caveat composes with

Algorithm 1 Directional audit: decide whether the predictor applies at all; if it does, locate the crossover and decide whether to add a floor

Require: deployed model h , held-out (X, y, a) of the vintage to be deployed on, fairness constraint C

Ensure: predicted direction, and whether a selection-rate floor is needed

```

1: if  $C$  does not constrain selection rates then
2:   return OUT OF SCOPE {e.g. equalized odds}
3: end if
4:  $\pi \leftarrow \max(\bar{y}, 1 - \bar{y})$  {trivial-predictor accuracy}
5:  $s \leftarrow h(X)$ 
6:  $\mathcal{V} \leftarrow \{\tau : \text{acc}(s > \tau) \geq \pi \wedge 0.05 \leq \overline{(s > \tau)} \leq 0.95\}$ 
7: if  $\text{span}(\mathcal{V}) < 0.40$  then
8:   return UNSWEEPABLE; compare natural sub-
     populations
9: end if
10:  $T \leftarrow 12$  thresholds from  $\mathcal{V}$ , spread evenly in selection rate
11: for all  $\tau \in T$  do
12:    $r_\tau \leftarrow \overline{(s > \tau)}$ ; fit  $C$  on the arm thresholded at  $\tau$ 
13:    $\Delta_\tau \leftarrow \%$  change in favourable decisions
14:   discard  $\tau$  if parity gap  $< 0.05$  or  $\text{acc} < \pi$ ; mark  $\tau$ 
     advisory if  $r_\tau < 0.10$  {route-divergence region}
15: end for
16: if fewer than 4 non-advisory arms remain, or  $\max \Delta - \min \Delta < 2$  points then
17:   return VOID {nothing moved; not a refutation}
18: end if
19: probe each fitted arm’s mixture for a flat keep-probability
     (Sec. VIII-F); flag any flat arm
20: if  $\text{sign}(\Delta_\tau)$  changes more than once across the non-
     advisory arms, or changes from  $+$  to  $-$  as  $r_\tau$  rises then
21:   return NON-MONOTONE {the rule does not apply
     here}
22: end if
23: if every non-advisory  $\Delta_\tau$  shares one sign then
24:   return that sign’s verdict {it holds across the whole
     viable band}
25: end if
26:  $c \leftarrow (\max\{r_\tau : \Delta_\tau < 0\}, \min\{r_\tau : \Delta_\tau > 0\})$ 
27:  $r_0 \leftarrow$  selection rate of the deployed model
28: if  $\min c < r_0 < \max c$ , or  $|\Delta|$  at the arms nearest  $r_0$  is
     within seed noise then
29:   return INDETERMINATE {the sealed record’s misses
     live here}
30: end if
31: if  $r_0 \leq \min c$  then
32:   return WITHDRAWAL; add a selection-rate floor
33: else
34:   return EXTENSION; the floor buys little
35: end if

```

Section VIII-F: these are natural arms, where the mixture is graded; a flat-lottery arm has no score-respecting extraction, and derandomizing it undoes the parity it bought.

TABLE VI

CROSSOVER LOCATIONS, WITH THE EXISTENCE OF A CROSSING SEPARATED FROM ITS LOCATION. **P(CROSSING)** IS THE SHARE OF NESTED ROWS-BY-SEEDS RESAMPLES IN WHICH A SIGN CHANGE CAN BE BRACKETED AT ALL; **CROSSOVER** IS THE MEDIAN LOCATION *conditional* ON ONE EXISTING. SEPARATING THEM IS NOT PEDANTRY: READ UNDER SEEDS ONLY, THE DUTCH CENSUS APPEARED TO CROSS AT 0.576 WITH A 95% INTERVAL OF 0.572–0.580 — AN APPARENTLY PRECISE ESTIMATE — WHILE UNDER THE NESTED BOOTSTRAP *no* RESAMPLE BRACKETS A CROSSING, AND THE PUBLISHED FIGURE WAS A PROPERTY OF THE SIX ARMS THAT PRODUCED IT. AN INTERVAL CONDITIONED ON AN EVENT THAT MAY NOT OCCUR CANNOT BE READ AS AN INTERVAL. THE BRACKET-ONLY ROWS ASSUME A CROSSING RATHER THAN ESTIMATING WHETHER THERE IS ONE, AND THE LENDING ROWS ARE POINT ESTIMATES WITH NO RESAMPLING; BOTH ARE LABELLED AS SUCH RATHER THAN SHARING A COLUMN WITH THE NESTED FIGURES. THE RATES ARE NOT ONE ECONOMIC OBJECT: SURVEY ROWS ARE PREDICTED-STATUS RATES OVER SAMPLED PERSONS, LENDING ROWS ARE APPROVAL RATES CONDITIONAL ON HAVING APPLIED, AND COMPAS INVERTS A RISK SCORE.

Population	Domain	P(crossing)	Crossover given one	Uncertainty
<i>Nested rows-by-seeds bootstrap, 1,000 pairs</i>				
COMPAS	crim. justice	100%	0.457	0.433–0.539
S. Carolina	income	97%	0.525	0.454–0.537
Oregon	income	99%	0.532	0.514–0.604
Dutch census	occupation	0%	—	none crosses
<i>Bracket only — crossing assumed, not estimated</i>				
Florida	income	—	0.284	0.259–0.309
Ohio	income	—	0.556	0.528–0.585
Pennsylvania	income	—	0.652	0.624–0.681
<i>Lending — one market, three overlapping estimates</i>				
HMDA pooled	mortgages	—	0.660	point only
HMDA Louisiana	mortgages	—	0.659	point only
HMDA refinance	mortgages	—	0.758	point only

What the audit costs. It fits the constraint $12 \times$ seeds times; with five seeds and the linear learners used here, one fit takes 2–10 seconds at 15k–69k training rows, so a full audit is minutes per population, scaling linearly in the base learner’s own fit time since each reduction call is a sequence of learner refits. On a production gradient-boosted model over millions of rows that multiplier is the binding cost, and we have not measured it there.

VIII. BOUNDARIES

A. Where the crossover sits, and what it cost to find out

The crossover was first published as a cluster, on four populations of one instrument. It did not survive measurement. **Located crossovers span 0.22 to 0.81** — 0.22 to 0.65 across the ACS income sweeps, with the law-school and South Carolina mortgage brackets above 0.80 — so the location is population-specific and only sometimes near the middle. Nothing in this work predicts it: the sealed test of the two candidate predictors returned UNDERPOWERED with three crossovers located against a floor of six, and three of ten large states have no crossover to locate at all, sweeping U-shaped or levelling up throughout. The lending estimates come from one two-state market read three overlapping ways, so “lending is different” remains a hypothesis rather than a finding.

That history is the case for measuring rather than transporting, and it is made by the prior’s own failures rather than

TABLE VII

SENSITIVITY TO THE PARITY-GAP EXCLUSION. CORRELATION PER POPULATION AS THE THRESHOLD IS LOOSENED.

Population	0.01	0.02	0.05	0.10
Dutch census	+0.851	+0.908	+0.946	+0.938
COMPAS	+0.844	+0.844	+0.844	+0.871
S. Carolina	+0.012	+0.012	+0.905	+0.849
Oregon	+0.095	+0.095	+0.664	—
Alabama	−0.368	−0.368	−0.368	+0.296
Kentucky	−0.590	−0.590	−0.654	—

against them. Sweeping the arms where a fixed 0.54 fails shows they are populations with no directional rule to get right — non-monotone or outright inverted, including one the rule scored *correct*, whose correctness was therefore luck. A prior read off a dashboard cannot detect that, because detecting it requires the sweep the prior exists to avoid. The audit refuses all of them; only the prior is exposed.

B. The procedure is unusable on very generous tasks

Varying the selection rate by moving a threshold works by *degrading* the classifier. On a task where most people already qualify, reaching a low selection rate requires a model that refuses qualified people, and past a point that model is worse than a constant. Defining the *viable band* as the range of selection rates reachable by a classifier still beating $\max(p, 1 - p)$: Dutch spans 0.895, COMPAS 0.859, typical ACS ≈ 0.51 , and LSAC only **0.056**. On LSAC, where 89% pass the bar, no choice of thresholds produces a usable sweep. This is computable before running anything.

When the viable band forbids a sweep, the crossover can still be located by comparing natural sub-populations — which is how the mortgage crossover was first found.

C. A sealed prediction, and what it cost

Every result above is a correlation computed after the arms were run, which is the weakest support for a claim of the form *you can know in advance*. We therefore ran the strongest format available: nine populations whose direction had never been measured, with the rule, the populations and the pass mark committed before any of them existed. Eight enter the sealed score; the ninth, Taiwan, was sealed alongside them but registered as reported separately, because a single non-Western population can carry no claim on its own (it was predicted correctly: rate 0.885, observed +0.68%).

It failed. Four of eight, against a bar of seven, not beating a constant (Table VIII).

What failed is a refinement, not the rule. Table VII shows the exclusion threshold doing real work: South Carolina moves from +0.905 to +0.012 when it is loosened. Investigating that, we found the newly admitted arms are not noise — at twenty seeds they are consistently *positive*, +8.73, +6.85, +9.09, six standard errors from zero — and concluded the relationship turns back up below a selection rate of 0.10. The sealed rule carried that refinement.

TABLE VIII

THE SEALED PREDICTION. POPULATIONS NEVER PREVIOUSLY MEASURED; PREDICTIONS COMMITTED BEFORE THE RUNS.

Population	Rate	Δ pool	Pred.	Actual
MO, \$100k	0.026	−16.48%	up	down
AZ, \$100k	0.037	−23.34%	up	down
IN, \$70k	0.081	−23.32%	up	down
TN, \$50k	0.255	−1.68%	down	down
MD, \$50k	0.508	+0.78%	down	up
MI, \$30k	0.612	+0.85%	up	up
NC, \$20k	0.789	+0.16%	up	up
GA, \$10k	0.905	+0.30%	up	up
Taiwan 2005	0.885	+0.68%	up	up

TABLE IX

THE RE-SEALED PREDICTION: THE UNREFINED RULE — DOWN BELOW 0.54, UP AT OR ABOVE — COMMITTED BEFORE TEN NEVER-MEASURED STATES RAN. BAR: 9 OF 10 *and* BEAT THE BEST CONSTANT.

Population	Rate	Δ pool	Pred.	Actual
NE, \$100k	0.014	−6.35%	down	down
AR, \$50k	0.195	−2.55%	down	down
OK, \$50k	0.220	−12.78%	down	down
WA, \$70k	0.233	−5.52%	down	down
NV, \$50k	0.297	−1.57%	down	down
MN, \$30k	0.699	−0.04%	up	down
CO, \$30k	0.701	+0.63%	up	up
KS, \$20k	0.767	+0.58%	up	up
WI, \$20k	0.814	+0.69%	up	up
IA, \$10k	0.896	+0.04%	up	up

All three held-out low-rate arms came out strongly *negative*. Scoring the same eight arms under the simpler rule we held before the refinement — down below the crossover, up above it — gives **seven of eight**, its only miss being Maryland at 0.508, within 0.032 of the crossover on an effect of +0.78%. That seven is post-hoc and we do not claim it; what was sealed scored four.

The refinement was route-specific. Its arms reach a rate near 0.05 by thresholding one fixed score vector; the sealed arms reach it with a genuinely rarer label. At that rate the two disagree completely, +8.73, +6.85, +9.09 against −16.48, −23.34, −23.32. So Section V-C qualifies: the two routes agree through the middle of the range and diverge at the bottom, and the accuracy rule does not catch the divergent arms — they clear the trivial-predictor bar and remain artifacts.

A refinement drawn from four in-sample populations, internally consistent and backed by a twenty-seed replication, made the prediction worse than the rule it replaced and no better than a constant. Nothing in the in-sample data could have shown that.

D. The rule re-sealed, and it holds

The unrefined rule — down below 0.54, up at or above — was committed with its ten populations and its pass mark before any of their arms existed: at least 9 of 10 correct *and* strictly beating the best constant. **It holds: 9 of 10, best constant 6** (Table IX). The re-seal applies none of this paper's tuned guards and scores all ten populations raw, so it has

nowhere to hide an exclusion, which makes it the cleanest scoring here.

Two statistics, because they answer different questions and the weaker one is the honest headline. A guesser at the constant's rate reaches 9 or more with probability 0.046. But rule and constant disagree on only five arms, the rule wins four of those five, and the one-sided sign test on the discordant calls gives $p \approx 0.19$. Nor is this the first sealed attempt at a direction rule, so the binomial tail carries a sequential correction whose size is a judgement we should not make in our own favour: read narrowly it is a factor of two and the tail is near 0.09; read as this paper's own abstract counts them — nine sealed or registered direction tests — it is a factor of nine and the tail is 0.41. Across the conventional line either way, and not close to it on the wider reading.

What the pass cannot establish, stated where the score is.

Every arm reaches its rate by the income-cutoff route, and the achieved rates leave the interior of (0.30, 0.70) unsampled — the nearest sit at 0.297, 0.699 and 0.701. So the pass cannot separate the selection rate from label rarity as the operative variable, nor the value 0.54 from any prior in that interval: **a rule reading only the income cutoff scores identically here.** The screen-gated race cohort was registered to break exactly this confound and did, on mixed-route arms off-instrument and off-continent, where the cutoff-only null scored 5 of 10 against the rate rules' 8–9 — ten arms from two person-samples, so the confound breaks within populations rather than across them. Its sealed S1 nonetheless *fails* by its own letter, because a 0.50 prior edged the 0.54 prior 9 to 8: the rate carries the prediction and the specific value does not.

The paired statistic has a ceiling this design cannot raise.

A third cohort was built to supply the discordant arms the sign test needs, and its pre-committed magnitude guard removed ten of twenty-four — removing them *near the crossover*, because effect size grows with distance from it ($r = +0.648$). Four discordant arms survived, the rule won all four, and four of four is $p = 0.0625$: the bar was arithmetically unreachable. The guard and the discordance a paired test needs are anti-correlated by construction, so the weakness does not yield to a larger balanced cohort. On a second, post-hoc cohort of larger states the fixed prior scored 5 of 10, and every miss agrees with that state's own measured crossover — which is the case for measuring rather than transporting, made by the prior's failures.

E. Small populations

Below roughly 2,500 test subjects the method's own randomness exceeds the entire effect of the constraint. This is a specific instance of predictive multiplicity [21], [25] — and Long *et al.* [21] show fairness interventions themselves exacerbate it — models with equivalent aggregate performance disagreeing substantially on individuals — and we position it as a caution supported by that literature rather than as a novel finding. COMPAS has 1,584 and flips sign between seeds on two of four arms; LSAC has 6,240 and flips on none. COMPAS carries direction here and never magnitude.

F. How the withdrawal is implemented: a lottery

At severely cut operating points the fitted mixture stops being a graded adjustment and becomes a lottery: nobody below the threshold receives any probability, and every predicted positive of both groups is kept at one flat rate — 0.65 on the Dutch census arm, 0.135 on COMPAS — with keep probability uncorrelated with the person's own score. The gap closes by random discard and the certificate is unchanged, which matters because a model that cannot generate a truthful individualized reason for a denial sits badly with adverse-action and notice duties. **It is a checkable hazard rather than a cost of parity:** a control at each population's natural operating point finds every mixture graded, so the signature belongs to constraints pushed to extremes and not to withdrawal as such — which is why Algorithm 1 probes for it (line 19) and why the full treatment, including the search of a richer two-threshold class that returns an empty feasible set, is reported separately.

IX. METHOD, AND WHAT WE GOT WRONG

Table X is the complete ledger — every registered or sealed test, with its committed bar and its outcome, failures included and uncorrected. Counted once and consistently it holds **eighteen** rows: twelve fail, three record a hold (one alongside a failed first stage), two returned underpowered, and one settled a reading rather than a pass or a fail. Of the nine that test a *direction* rule on populations it had not seen, one passed. The rows should not be read as eighteen attempts at one claim of which three succeeded: each targeted a different claim, each failure removed a stronger claim than the one that survives, and the last column records what each row settled. The failures are boundary measurements; the passes are tests of the claim *as those boundaries define it*.

Three episodes are worth stating in the body because each cost a claim the paper would otherwise still be making. The rule does not survive equalized odds ($r = +0.334$ against a $+0.70$ bar), which is why the claim reads *parity* and not *group fairness*: parity controls approval rates directly and equalized odds does not. Our attempt to *derive* the rule was beaten by a constant, which is why the paper is empirical and says *predicts* rather than *causes*. And a refinement drawn from four in-sample populations, internally consistent and backed by a twenty-seed replication, was sealed, scored 4 of 8, and made the prediction worse than the unrefined rule it replaced — nothing in the in-sample data could have shown that, which is the whole argument for sealing.

Two published results of our own are withdrawn here rather than quietly corrected: two correlations previously reported at $+0.979$ and $+0.968$ are void, exposed by an exchange rate of exactly 0.000 — the post-processor had returned an identical model at every operating point, so the effect was a constant divided by a shrinking baseline. A published crossover was also downgraded by a bootstrap we chose to run against our own result.

TABLE X

THE EIGHT REGISTERED TESTS THAT IMPOSE A SCOPE CONDITION ON THE SURVIVING CLAIM. THE LEDGER RUNS TO EIGHTEEN ROWS — TWELVE FAIL, THREE HOLD, TWO UNDERPOWERED, ONE A READING RATHER THAN A VERDICT — AND THE COMPLETE RECORD CARRIES ALL OF THEM WITH THEIR REGISTRATION AND SCORING COMMITS. NOTHING IS WITHDRAWN HERE; TEN ROWS ARE ELSEWHERE.

Test	Bar / outcome	What it settled
EO transfer	$r \geq +0.70$; $+0.334$ — fail	the rule works for parity, which controls approval rates directly; it failed for equalized odds, which does not
Post-processing transfer	carry over; $r = -0.024$ — fail	the rule stops where the model may read the protected attribute — the boundary of Sec. VI
Derivation	beat a constant; beaten — fail	our own attempt to derive the rule lost to a constant; the first-order account we could not extend to the blind randomized case is named in Sec. II
Sealed rule, refined	7 of 8; got 4 — fail	the extra clause added from early data was wrong, and is withdrawn
Re-seal, unrefined rule	9 of 10 + beat constant; got 9, constant 6 — holds	the simple rule predicts populations it has never seen, from the rate alone
Sealed magnitude model	beat predict-zero; MAE 4.50 v. 0.77 — fail	size cannot be predicted from distance and band span; the direction-only concession is earned
Regime deconfounding	two sealed readings; R1 18/18, $r = +0.672$ — method reading	the pool boundary is the optimizer family, not blindness; the theorem's group directions hold in every aware cell (Sec. VI)
Race cohort (screened)	S1: 8/10, but 0.5-null 9 — fail; S2 — holds	the cutoff null is beaten 8–9 vs 5: it is the rate , not label rarity; the 0.54 value holds no in-band privilege; within-population passes off-instrument at $\rho \geq 0.9$

X. LIMITATIONS

Most of the evidence is about predicted labels, not allocations. In five of the eight sources no one receives anything: “favourable decisions” are a classifier’s labels over people whose incomes, occupations or bar results already exist. HMDA is the one source recording real allocative decisions — and it is also where the direction reverses and where our sweep protocol failed its held-out test, so whether the construct transports from status-prediction benchmarks to allocation data is an open question rather than an assumption we are entitled to. The practical advice of Section VII is addressed to teams making real allocations, and they should read this paragraph first.

The lending evidence is thinner than its market count suggests. Forty markets carry a rate-bearing arm on each attribute, but every one of the twenty-six sex arms falls below the audit’s own disparity floor — median baseline gap 0.013 against the race arms’ 0.079 — so the allocation result rests entirely on race arms. That is a property of US mortgage approval rather than a gap in the method, and the audit refuses those arms correctly. Nor does either lending count discriminate: every arm sits at a rate of 0.82 or above, so the rule predicts extension on all of them and a constant scores identically. US mortgage approval supplies no market whose

rate straddles a crossover, so it cannot test the rule at all.

Magnitude does not transfer between populations, only within them, and a sealed model of it lost to predicting zero. **The crossover’s location is unexplained and wider than it looked:** located values span 0.22 to 0.81, the sealed test of the two candidate predictors returned underpowered, and some populations have no crossover to locate. **We cannot say why the rule works.** Our own derivation lost to a constant; the first-order account we could not extend to the blind randomized case is named in Section II.

A label is not a task. A fixed nominal outcome definition is a different task each vintage: a dollar cutoff slides down an inflated income distribution, and restoring the earlier shape requires matching the label’s *quantile*, not its real price — CPI-anchoring is not enough. Four 2022 populations invert for this reason, and one 2014 state inverts at a lower cutoff, so this is a property of nominal labels rather than of a vintage.

Coverage is eight data sources, five countries and two decades, of which 152 of 161 populations come from two instruments; the eight sources remain the honest count of independent instruments. Taiwan, Brazil and Mexico are the non-Western populations and arrived late enough that most of the record predates them. The only EU population is the 2001 Dutch census, which by this paper’s own vintage standard grounds no claim about current EU deployments.

XI. DISCUSSION

What a practitioner can use. Run the audit on your own data. It will locate your crossover, tell you your system admits no directional rule, or tell you the question is moot because the pool barely moves across your whole viable band — and each of those is an answer. Where it returns a direction and that direction is withdrawal, a selection-rate floor is the remedy: not a new method, but a linear moment constraint inside the reductions framework [1] and a variant of the minimum-rate constraints of [11]. It costs about 0.12 accuracy points and moves the exchange rate from 1.47 destroyed per created to 0.88. Two honest consequences follow. Because the floor is that cheap, a team that expects withdrawal and does not want to sweep should simply apply it; what the audit adds is the ability to tell in advance whether it is needed, how far the operating point sits from the sign change, and whether the population admits an answer at all. And whatever the audit returns, the deployed system’s pool change and who-pays decomposition should be measured afterwards and reported beside the parity number — so that escaping the ex-ante audit buys no escape from the accounting.

What to report. The parity difference alone certifies both outcomes this paper distinguishes. A report that adds the pool change and the group-conditional flows costs nothing to compute, requires no doctrinal commitment, and is the difference between a dashboard that can tell those outcomes apart and one that cannot.

Adversarial inputs. A procedure that can refuse can also be gamed, in four places we can name: the label definition, the segmentation the audit is run over, the sweep’s own

parameters, and refusal used as an exit. The countermeasures are symmetrical — fix the label’s quantile rather than its nominal value, pre-register the segmentation, freeze the sweep grid, and pair every verdict including a refusal with the ex-post accounting above. One concrete instance is worth stating: on HMDA a single constraint earns “extension” on the pooled book while concentrating withdrawal in home improvement, so a segmentation chosen after the fact can hide the thing the audit exists to surface.

What we are not claiming. The rule is not surprising once stated — say yes often and the constraint tends to give, say yes rarely and it tends to take. What is not available without measuring is where the boundary falls for a particular population, and located crossovers span 0.22 to 0.81. In the one setting where we drew populations at random rather than choosing them, no population’s operating rate reached 0.55 and half of those with a computable rate sat inside the band where an unaided reading of “often” against “rarely” has nothing to say. A further fifth follow no directional rule at all, and there the intuition returns a confident answer while the audit declines to.

APPENDIX: TWO WORKED TRACES OF ALGORITHM 1

Oregon 2018, sex, to a verdict. Demographic parity constrains selection rates: in scope (line 1). Trivial-predictor accuracy $\pi = 0.636$ (line 4); the viable band clears the 0.40 span bar, so the sweep proceeds (lines 6–9). The stored sweep holds 18 arms; the guards discard 2 on the parity-gap floor and 2 as worse than doing nothing, leaving 14 spanning rates 0.111–0.653, none below the 0.10 advisory boundary (lines 11–15). Their pool changes span 7.14 points, clearing the void guard (line 16). The signs, sorted by rate, are -----+++: one rising change, so the landscape is monotone-crossing (lines 20–25) and the bracket is $c = (0.528, 0.589)$ (line 26). The deployed model’s rate is $r_0 = 0.353$ — below $\min c$, outside seed noise: verdict WITHDRAWAL, add a floor (line 32). What the constraint in fact does at Oregon’s natural arm: –3.0% of the pool.

LSAC bar passage, to a refusal. Trivial-predictor accuracy $\pi = 0.890$ — almost everyone passes. The viable band spans 0.056, far under 0.40: verdict UNSWEEPABLE at line 8, before a single constraint is fitted. Had the guard been ignored, the stored sweep shows what awaited: five of seven arms fall below the trivial predictor and the two survivors are too few to read — VOID. Both endings refuse; neither manufactures a direction.

REPRODUCIBILITY

Code, stored results and the full document trail are in the accompanying repository, which holds the per-test registration and scoring commits, the OpenTimestamps receipts for the seals that carry them, and the complete eighteen-row ledger. Every figure here is re-derived from stored results by a test suite that fails if a documented number and its experiment disagree, and counts appearing in more than one place are derived from one source and checked against each other —

after an audit found four descriptions of the same table. One limit cannot be repaired retroactively: for ten of the sixteen seals the ordering rests on git commit timestamps from our own machine, which an outsider must take on trust. The later seals carry external anchoring; the earlier ones do not.

REFERENCES

- [1] A. Agarwal, A. Beygelzimer, M. Dudík, J. Langford, and H. Wallach, “A reductions approach to fair classification,” in *Proc. ICML*, 2018.
- [2] B. H. Zhang, B. Lemoine, and M. Mitchell, “Mitigating unwanted biases with adversarial learning,” in *Proc. AAAI/ACM Conf. AI, Ethics, and Society (AIES)*, 2018, pp. 335–340.
- [3] T. Kamishima, S. Akaho, H. Asoh, and J. Sakuma, “Fairness-aware classifier with prejudice remover regularizer,” in *Machine Learning and Knowledge Discovery in Databases (ECML PKDD)*, LNCS 7524, Part II, 2012, pp. 35–50.
- [4] M. Feldman, S. A. Friedler, J. Moeller, C. Scheidegger, and S. Venkatasubramanian, “Certifying and removing disparate impact,” in *Proc. ACM SIGKDD*, 2015, pp. 259–268.
- [5] M. Hardt, E. Price, and N. Srebro, “Equality of opportunity in supervised learning,” in *Advances in Neural Information Processing Systems* 29, 2016.
- [6] L. Hu and Y. Chen, “Fair classification and social welfare,” in *Proc. ACM Conf. Fairness, Accountability, and Transparency (FAT*)*, 2020, pp. 535–545.
- [7] L. T. Liu, S. Dean, E. Rolf, M. Simchowitz, and M. Hardt, “Delayed impact of fair machine learning,” in *Proc. ICML*, PMLR 80, 2018, pp. 3150–3158.
- [8] A. K. Menon and R. C. Williamson, “The cost of fairness in binary classification,” in *Proc. Conf. Fairness, Accountability and Transparency*, PMLR 81, 2018, pp. 107–118.
- [9] S. Corbett-Davies, J. D. Gaebler, H. Nilforoshan, R. Shroff, and S. Goel, “The measure and mismeasure of fairness,” *J. Mach. Learn. Res.*, vol. 24, no. 312, pp. 1–117, 2023.
- [10] X. Zeng, E. Dobriban, and G. Cheng, “Bayes-optimal classifiers under group fairness,” arXiv:2202.09724, 2022.
- [11] B. Mittelstadt, S. Wachter, and C. Russell, “The unfairness of fair machine learning: Levelling down and strict egalitarianism by default,” *Michigan Technology Law Review*, vol. 30, no. 1, 2024.
- [12] N. Krco, T. Laugel, V. Grari, J.-M. Loubes, and M. Detynecki, “When mitigating bias is unfair: Multiplicity and arbitrariness in algorithmic group fairness,” arXiv:2302.07185, 2023.
- [13] G. Maheshwari, A. Bellet, P. Denis, and M. Keller, “Fair without leveling down: A new intersectional fairness definition,” in *Proc. EMNLP*, 2023.
- [14] Y. Yang, X. Chang, and P.-y. Chen, “Fairness may backfire: When leveling-down occurs in fair machine learning,” arXiv:2603.06901, 2026.
- [15] M. Kearns, S. Neel, A. Roth, and Z. S. Wu, “Preventing fairness gerrymandering: Auditing and learning for subgroup fairness,” in *Proc. ICML*, 2018.
- [16] F. Ding, M. Hardt, J. Miller, and L. Schmidt, “Retiring Adult: New datasets for fair machine learning,” in *Proc. NeurIPS*, 2021.
- [17] S. Goethals, E. Delaney, B. Mittelstadt, and C. Russell, “Resource-constrained fairness,” arXiv:2406.01290, 2024.
- [18] J. Angwin, J. Larson, S. Mattu, and L. Kirchner, “Machine bias,” *ProPublica*, 2016.
- [19] S. Corbett-Davies, E. Pierson, A. Feller, S. Goel, and A. Huq, “Algorithmic decision making and the cost of fairness,” in *Proc. KDD*, 2017.
- [20] E. Diana, W. Gill, M. Kearns, K. Kenthapadi, and A. Roth, “Minimax group fairness: Algorithms and experiments,” in *Proc. AAAI/ACM AIES*, 2021.
- [21] C. X. Long, H. Hsu, W. Alghamdi, and F. P. Calmon, “Individual arbitrariness and group fairness,” in *Proc. NeurIPS*, 2023.
- [22] S. Agarwal and A. Deshpande, “On the power of randomization in fair classification and representation,” in *Proc. FAccT*, 2022.
- [23] A. Cotter, M. Gupta, and H. Narasimhan, “On making stochastic classifiers deterministic,” in *Proc. NeurIPS*, 2019.
- [24] N. Grgić-Hlača, M. B. Zafar, K. P. Gummadi, and A. Weller, “On fairness, diversity and randomness in algorithmic decision making,” arXiv:1706.10208, 2017.
- [25] E. Black, M. Raghavan, and S. Barocas, “Model multiplicity: Opportunities, concerns, and solutions,” in *Proc. FAccT*, 2022.

- [26] S. Barocas and A. D. Selbst, “Big data’s disparate impact,” *California Law Review*, vol. 104, no. 3, pp. 671–732, 2016.
- [27] *Ricci v. DeStefano*, 557 U.S. 557 (2009).
- [28] J. Broome, “Fairness,” *Proceedings of the Aristotelian Society*, vol. 91, no. 1, pp. 87–102, 1990–91.
- [29] P. Stone, *The Luck of the Draw: The Role of Lotteries in Decision Making*. Oxford University Press, 2011.