

The Decision I Need

SUBThe Decision I Need

Dheirav

August 13, 2026

The decision I need

The question

Is the scope condition enough on its own for FAccT or AIES, or should I establish the mechanism first?

Why it is a real fork

If the scope condition is enough, the paper is close to submittable. The claim — levelling down reverses above a selection rate of roughly 0.3 — is established by a single-factor design, replicated across two protected attributes, five states, a second domain, and fourteen held-out populations. The remaining work is writing, figures, and one missing baseline.

If the mechanism is needed first, that is a different project. My derivation cleared its pre-registered bars and was then beaten by a constant rule, so I have no explanation. Making one would require populations where curvature and selection rate come apart — score distributions with the same operating rate and different shapes — which is a harder experiment than anything run so far.

What I would argue

The scope condition is the contribution, and it stands on its own for three reasons:

1. It **corrects a widely-cited claim**. Mittelstadt et al. put "by default" in their title; the direction is conditional, and they were working inside the one regime where their framing holds.
2. It is **causally identified**, not correlational. One factor moves and the test code asserts the rest is fixed.
3. It has a **policy reading that does not need a mechanism**: lending, hiring and admissions are low-selection-rate settings, so the regime where these constraints take decisions away is the regime they are deployed in.

Plenty of accepted empirical work establishes *when* something happens without establishing *why*.

The counter-argument I expect

A reviewer may say a scope condition without a mechanism is a curve, and that the crossover value (0.25–0.60) is too imprecise to act on. That is fair. My answer is that the direction is what changes

a deployment decision, not the third decimal place of the threshold — but I would rather hear whether that holds up before writing the paper around it.

Secondary questions, if there is time

1. **Venue and timing.** FAccT or AIES? The deadline sets the page limit, which decides what gets cut from 16 research documents.
2. **The missing baseline.** Mittelstadt et al.’s construction is a per-group floor applied by post-processing. I have never run it directly — the closest arm uses group-wise thresholds, which is not the same thing. Worth an hour. Should I?
3. **Authorship and AI assistance.** I used an AI assistant substantially for implementation, analysis and drafting. What does the department expect me to disclose, and how?

What is not in question

The empirical work is verified. Every headline number is re-derived from stored results by an automated check that fails if a document and its data disagree. Three of my own pre-registered predictions failed and are reported as failures. Two of my own claims were retracted in place. I would rather be told the framing is wrong now than discover it in review — which is exactly what happened with the remedy, and why it is no longer claimed as novel.