The Decision I Need

One fork, the argument for it, and the counter-argument expected

Dheirav

August 25, 2026

The decision I need

Status, one line: 67 populations, 8 data sources, 5 countries, sixteen registered tests with ten failures published uncorrected, an 18-page draft, and the one standing confound resolved this week in the finding's favour. The full picture is in the handout and the status report; this document argues the single question I most need answered.

The question

A March 2026 theory paper proves the effect I study can go either way, with no experiments in it. I have the experiments, a rule a practitioner can apply, and results that say the theory is right about one half of its claim while its boundary needed correcting on the other. **Is that an empirical paper worth publishing, or does it read as a follow-up to someone else's result?**

Why it is a genuine fork

If the empirical half stands on its own, the paper is close to ready. Their paper proves the possibility and contains zero experiments. Mine has 67 populations across six kinds of decision, a controlled design that separates the approval rate from the difficulty of the task, a sealed prediction record, and a rule computable from a number every deployed system already has.

The strongest thing I have is a two-sided result about their theory:

- Their conditions, the orderings their theorem is stated over, could not be evaluated on any of 26 real datasets under either of two estimators I tried. They are sufficient rather than necessary, so this says their theorem is silent on real data, not that it is wrong. A relaxed form of the same ordering tracks the observed direction almost everywhere, and my approval rate tracks their quantity at 0.85 across 150 populations.
- Their group-level prediction, that the disadvantaged group gains wherever the attribute is readable, held in every one of the 45 populations I measured that way. My sealed deconfounding test also corrected where their regime boundary bites in practice: it is the kind of method, not attribute blindness.

So I am not applying their theory and not refuting it. I am showing which half survives contact with data, correcting where its line falls, and supplying the instrument it lacks. That is what an empirical paper can settle and a proof cannot.

If it instead reads as derivative, that is a framing problem rather than a content problem, but it is still a problem. I reached the finding before knowing their paper existed, which is provable from dates, but a reader encountering both will see theirs first.

What I would argue

1. Theory says it can go either way. Only measurement says when, and the "when" is what a practitioner is standing in.
2. Their conditions do not work on data; mine needs last year's approval rate. That difference is the contribution.
3. The direction question the whole record kept asking, whether it is truly the rate or an artifact of how I varied it, was answered this week by a sealed test on Brazilian census data: the artifact reading scored 5 of 10, the rate scored 8.
4. The practical claim is sharper than it sounds: two products in the same mortgage market, run through the same fairness constraint, move in opposite directions, and the fairness report says success in both cases.

The objection I expect

That this is the empirical appendix to someone else's theory paper.

My answer: their result cannot be applied to data at all, so this is not an application of it. It is a demonstration that it cannot be applied, plus a substitute that can, plus a correction their side must absorb. But I would rather be told now if that distinction is too fine to carry a paper.

What is not in question

The measurements. Every headline number is regenerated from stored results by automated checks that fail if a document and its data disagree, and every sealed prediction's timing is verifiable from the public repository, the recent ones anchored in the Bitcoin blockchain.

What I would want you to press me on

I withdrew two headline results when a guard I added showed their arms were worse than doing nothing, and my strongest sealed pass carries its own health warning in print: paired against the best lucky constant it is suggestive rather than significant. The method that produced those admissions, registering every prediction with the naive alternative it must beat, is the part of this work I would most like judged, because it kept producing results I did not want and published them anyway.

The remaining smaller decisions, venue, title, and whether to contact the theory paper's authors, are listed on the handout's last page.