

The Decision I Need

One fork, the argument for it, and the counter-argument expected

Dheirav

August 25, 2026

The decision I need

The question. A March 2026 theory paper, "Fairness May Backfire" (arXiv:2603.06901), proves the effect I study can go either way, with no experiments in it. I have the experiments, a rule a practitioner can apply, and results that say the theory is right about one half of its claim while its boundary needed correcting on the other. **Is that a paper worth publishing, or does it read as a follow-up to someone else's result?**

Side by side.

	Fairness May Backfire (theirs)	This work (mine)
What it is	A proof; zero experiments, zero datasets	Measurement across 67 populations, 8 sources, 5 countries
The claim	The direction of a fairness fix depends on the data; it can go either way	Which way, readable in advance from the baseline approval rate against a measured crossover
Can it be used?	Its conditions could not be evaluated on any of 26 real datasets, under two different estimators	Needs one number every deployed system already has; the rate tracks their quantity at 0.85 across 150 populations
Validation	None possible without experiments	Sixteen registered predictions, pass marks committed before the data existed, ten failures published uncorrected, seals timestamped into Bitcoin
What each gives the other	The possibility result, and a group-level theorem my data confirms in all 45 aware populations	The instrument, the boundary map, and a correction: their regime line sits at the kind of method, not attribute blindness

The one standing doubt, whether it is truly the rate doing the predicting rather than an artifact of my test design, was resolved this week by a sealed test on Brazilian census data: the artifact reading scored 5 of 10, the rate scored 8. And the correction row matters for the question above, because appendices do not correct the papers they depend on.

The objection I expect. That this is the empirical appendix to their theory. My answer: a result that cannot be applied to data has not been applied by me. This is a demonstration that it cannot be applied, plus a substitute that can, plus a correction their side must absorb. I would rather hear now if that distinction is too fine.

What I would want you to press me on. I withdrew two headline results when my own guard exposed them, and my strongest sealed pass says "suggestive, not significant" in print. The method that forced those admissions is the part of this work I would most like judged.

The smaller decisions, venue, title, and contacting the theory authors, are on the handout's last page.